

Machine Learning

Lecture 2 Linear Models

Vincent Adam & Vicenç Gómez

2022-23

Content

- 1 Supervised learning
- 2 Perceptron
- 3 Linear regression
- 4 Logistic regression
- 5 Non-linear transforms
- 6 Exercises

Content

- 1 Supervised learning
- 2 Perceptron
- 3 Linear regression
- 4 Logistic regression
- 5 Non-linear transforms
- 6 Exercises

Supervised learning problem

A supervised learning problem consists of:


- A domain set $\mathcal{X} = \mathcal{X}^1 \times \dots \times \mathcal{X}^d$
- An **unknown** probability distribution \mathcal{D} on \mathcal{X}
- A target set \mathcal{Y}
- An **unknown** labelling function $f : \mathcal{X} \rightarrow \mathcal{Y}$
- A training set $S = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m))$ **sampled** from \mathcal{D} and f

Supervised learning

Given a supervised learning problem, the learner chooses the following:

- A hypothesis class \mathcal{H} of candidate labelling functions
- A loss function $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$
- An algorithm \mathcal{A} that minimizes the empirical risk

Linear models



- Simplest way to separate data is a **line** (or a **hyperplane** in higher dimensions)
- Hypothesis class: linear combination of features

Content

- 1 Supervised learning
- 2 Perceptron**
- 3 Linear regression
- 4 Logistic regression
- 5 Non-linear transforms
- 6 Exercises

Perceptron

- Algorithm for **binary classification**: $\mathcal{Y} = \{-1, +1\}$

Perceptron

- Algorithm for **binary classification**: $\mathcal{Y} = \{-1, +1\}$
- Assume inputs $\mathbf{x} = (x_1, \dots, x_d)$ on d **numerical** features

Perceptron

- Algorithm for **binary classification**: $\mathcal{Y} = \{-1, +1\}$
- Assume inputs $\mathbf{x} = (x_1, \dots, x_d)$ on d **numerical** features
- Compute a **weighted score** and output $+1$ or -1 as

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum_{i=1}^d w_i x_i > \theta \\ -1 & \text{otherwise} \end{cases}$$

Perceptron

- Sign function $\text{sign}(x)$ outputs $+1$ if $x > 0$ and -1 otherwise
- Hypothesis h completely defined by **weights** $\mathbf{w} = w_0, \dots, w_d$:

$$h(\mathbf{x}) = \text{sign} \left(\sum_{i=1}^d w_i x_i - \theta \right)$$

Perceptron

- Sign function $\text{sign}(x)$ outputs $+1$ if $x > 0$ and -1 otherwise
- Hypothesis h completely defined by **weights** $\mathbf{w} = w_0, \dots, w_d$:

$$h(\mathbf{x}) = \text{sign} \left(\sum_{i=1}^d w_i x_i - \theta \right) = \text{sign} \left(\sum_{i=0}^d w_i x_i \right) = \text{sign} (\mathbf{w}^\top \mathbf{x})$$

Perceptron

- Sign function $\text{sign}(x)$ outputs $+1$ if $x > 0$ and -1 otherwise
- Hypothesis h completely defined by **weights** $\mathbf{w} = w_0, \dots, w_d$:

$$h(\mathbf{x}) = \text{sign} \left(\sum_{i=1}^d w_i x_i - \theta \right) = \text{sign} \left(\sum_{i=0}^d w_i x_i \right) = \text{sign} (\mathbf{w}^\top \mathbf{x})$$

Perceptron

- Sign function $\text{sign}(x)$ outputs $+1$ if $x > 0$ and -1 otherwise
- Hypothesis h completely defined by **weights** $\mathbf{w} = w_0, \dots, w_d$:

$$h(\mathbf{x}) = \text{sign} \left(\sum_{i=1}^d w_i x_i - \theta \right) = \text{sign} \left(\sum_{i=0}^d w_i x_i \right) = \text{sign} (\mathbf{w}^T \mathbf{x})$$

where $x_0 = 1$ is a **dummy feature**
and $w_0 = -\theta$ (a.k.a. bias weight)

Perceptron

- Sign function $\text{sign}(x)$ outputs $+1$ if $x > 0$ and -1 otherwise
- Hypothesis h completely defined by **weights** $\mathbf{w} = w_0, \dots, w_d$:

$$h(\mathbf{x}) = \text{sign} \left(\sum_{i=1}^d w_i x_i - \theta \right) = \text{sign} \left(\sum_{i=0}^d w_i x_i \right) = \text{sign} (\mathbf{w}^T \mathbf{x})$$


where $x_0 = 1$ is a **dummy feature**
and $w_0 = -\theta$ (a.k.a. bias weight)

- Classification loss:

$$\ell(h(\mathbf{x}_i), y_i) = \mathbb{I}[h(\mathbf{x}_i) \neq y_i]$$

$$L_S(h) = \frac{1}{m} \sum_{i=1}^m \mathbb{I}[h(\mathbf{x}_i) \neq y_i]$$

Linearly separable data



Perceptron learning algorithm (PLA)

Builds a sequences of weights $\mathbf{w}^0, \mathbf{w}^1, \dots, \mathbf{w}^t$

Perceptron learning algorithm

- 1 Initialize weight vector $\mathbf{w}^0 = 0$
- 2 Find a **mistake** (\mathbf{x}_i, y_i) such that $h(\mathbf{x}_i) = \text{sign}(\mathbf{w}_0^\top \mathbf{x}_i) \neq y_i$
- 3 Update weights as $\mathbf{w}^1 \leftarrow \mathbf{w}^0 + y_i \mathbf{x}_i$
- 4 Repeat from 2. for weight vector \mathbf{w}^t , $t = 1, 2, \dots$

Properties


- If data is **linearly separable**, PLA guaranteed to converge to a hypothesis h_S (i.e. weight vector \mathbf{w}_S) such that $L_S(h_S) = 0$
- If data is **not** linearly separable, PLA never converges
- Variants:
 - Fix number of iterations T , stop when $t > T$
 - **Pocket algorithm**: only update weight vector \mathbf{w}^t when total number of mistakes decreases

Content

- 1 Supervised learning
- 2 Perceptron
- 3 Linear regression**
- 4 Logistic regression
- 5 Non-linear transforms
- 6 Exercises

Linear regression

- Assumes **regression problem** ($\mathcal{Y} = \mathbb{R}$)
- Hypothesis $h(\mathbf{x}) = \sum_{i=0}^d w_i x_i = \mathbf{w}^\top \mathbf{x}$
- Squared loss: $\ell(h(\mathbf{x}_i), y_i) = (h(\mathbf{x}_i) - y_i)^2$
- $L_S(h) = \frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}_i) - y_i)^2$ is the **mean squared error** (MSE)



Linear regression

- Hypothesis h defined by **weight vector** $\mathbf{w} = (w_0, \dots, w_d)$
- Find w that minimizes

$$L_S(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}_i) - y_i)^2 = \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2$$

- $L_S(\mathbf{w})$ is **continuous**, **differentiable**, and **convex**

Linear regression

$$\left\{ L_S(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{m} \left\| \begin{pmatrix} \mathbf{w}^\top \mathbf{x}_1 - y_1 \\ \vdots \\ \mathbf{w}^\top \mathbf{x}_m - y_m \end{pmatrix} \right\|^2 \right.$$

Linear regression

$$\left\{ \begin{aligned} L_S(\mathbf{w}) &= \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{m} \left\| \begin{pmatrix} \mathbf{w}^\top \mathbf{x}_1 - y_1 \\ \vdots \\ \mathbf{w}^\top \mathbf{x}_m - y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \left\| \begin{pmatrix} \text{---} & \mathbf{x}_1^\top & \text{---} \\ \text{---} & \mathbf{x}_2^\top & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_m^\top & \text{---} \end{pmatrix} \begin{pmatrix} w_0 \\ \vdots \\ w_d \end{pmatrix} - \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \right\|^2 \end{aligned} \right.$$

Linear regression

$$\left\{ \begin{aligned} L_S(\mathbf{w}) &= \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{m} \left\| \begin{pmatrix} \mathbf{w}^\top \mathbf{x}_1 - y_1 \\ \vdots \\ \mathbf{w}^\top \mathbf{x}_m - y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \left\| \begin{pmatrix} \text{---} & \mathbf{x}_1^\top & \text{---} \\ \text{---} & \mathbf{x}_2^\top & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_m^\top & \text{---} \end{pmatrix} \begin{pmatrix} w_0 \\ \vdots \\ w_d \end{pmatrix} - \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 \end{aligned} \right.$$

Linear regression

$$\left\{ \begin{aligned}
 L_S(\mathbf{w}) &= \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{m} \left\| \begin{pmatrix} \mathbf{w}^\top \mathbf{x}_1 - y_1 \\ \vdots \\ \mathbf{w}^\top \mathbf{x}_m - y_m \end{pmatrix} \right\|^2 \\
 &= \frac{1}{m} \left\| \begin{pmatrix} \text{---} & \mathbf{x}_1^\top & \text{---} \\ \text{---} & \mathbf{x}_2^\top & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_m^\top & \text{---} \end{pmatrix} \begin{pmatrix} w_0 \\ \vdots \\ w_d \end{pmatrix} - \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \right\|^2 \\
 &= \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = \frac{1}{m} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y})
 \end{aligned} \right.$$

Linear regression

$$\left\{ \begin{aligned} L_S(\mathbf{w}) &= \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{m} \left\| \begin{pmatrix} \mathbf{w}^\top \mathbf{x}_1 - y_1 \\ \vdots \\ \mathbf{w}^\top \mathbf{x}_m - y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \left\| \begin{pmatrix} \text{---} & \mathbf{x}_1^\top & \text{---} \\ \text{---} & \mathbf{x}_2^\top & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_m^\top & \text{---} \end{pmatrix} \begin{pmatrix} w_0 \\ \vdots \\ w_d \end{pmatrix} - \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = \frac{1}{m} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) \end{aligned} \right.$$

Linear regression

$$\left\{ \begin{aligned} L_S(\mathbf{w}) &= \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{m} \left\| \begin{pmatrix} \mathbf{w}^\top \mathbf{x}_1 - y_1 \\ \vdots \\ \mathbf{w}^\top \mathbf{x}_m - y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \left\| \begin{pmatrix} \text{---} & \mathbf{x}_1^\top & \text{---} \\ \text{---} & \mathbf{x}_2^\top & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_m^\top & \text{---} \end{pmatrix} \begin{pmatrix} w_0 \\ \vdots \\ w_d \end{pmatrix} - \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \right\|^2 \\ &= \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = \frac{1}{m} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) \end{aligned} \right.$$

- \mathbf{X} : $m \times (d + 1)$ matrix of inputs
- \mathbf{y} : $m \times 1$ vector of labels

Linear regression

- Minimize $L_S(\mathbf{w}) \Leftrightarrow$ set **gradient** $\nabla_{\mathbf{w}} L_S(\mathbf{w})$ to 0

$$\nabla_{\mathbf{w}} L_S(\mathbf{w}) = \begin{pmatrix} \frac{\partial L_S(\mathbf{w})}{\partial w_0} \\ \vdots \\ \frac{\partial L_S(\mathbf{w})}{\partial w_d} \end{pmatrix}$$

Linear regression

- Minimize $L_S(\mathbf{w}) \Leftrightarrow$ set **gradient** $\nabla_{\mathbf{w}} L_S(\mathbf{w})$ to 0

$$\nabla_{\mathbf{w}} L_S(\mathbf{w}) = \begin{pmatrix} \frac{\partial L_S(\mathbf{w})}{\partial w_0} \\ \vdots \\ \frac{\partial L_S(\mathbf{w})}{\partial w_d} \end{pmatrix}$$

- Gradient of loss term $(\mathbf{w}^\top \mathbf{x}_i - y_i)^2$ w.r.t. scalar weight w_k :

$$\frac{\partial (\mathbf{w}^\top \mathbf{x}_i - y_i)^2}{\partial w_k} = 2(\mathbf{w}^\top \mathbf{x}_i - y_i) x_{i,k}$$

Linear regression

- Minimize $L_S(\mathbf{w}) \Leftrightarrow$ set **gradient** $\nabla_{\mathbf{w}} L_S(\mathbf{w})$ to 0

$$\nabla_{\mathbf{w}} L_S(\mathbf{w}) = \begin{pmatrix} \frac{\partial L_S(\mathbf{w})}{\partial w_0} \\ \vdots \\ \frac{\partial L_S(\mathbf{w})}{\partial w_d} \end{pmatrix}$$

- Gradient of loss term $(\mathbf{w}^\top \mathbf{x}_i - y_i)^2$ w.r.t. scalar weight w_k :

$$\frac{\partial (\mathbf{w}^\top \mathbf{x}_i - y_i)^2}{\partial w_k} = 2(\mathbf{w}^\top \mathbf{x}_i - y_i) x_{i,k}$$

- **Gradient**

$$\nabla_{\mathbf{w}} L_S(\mathbf{w}) = \frac{2}{m} (\mathbf{X}^\top \mathbf{X} \mathbf{w} - \mathbf{X}^\top \mathbf{y})$$


Linear regression

- $\nabla_{\mathbf{w}} L_S(\mathbf{w}) = 0 \Leftrightarrow \mathbf{X}^\top \mathbf{X} \mathbf{w} = \mathbf{X}^\top \mathbf{y}$
- $\mathbf{X}^\top \mathbf{X}$ is a $(d+1) \times (d+1)$ matrix

Linear regression

- $\nabla_{\mathbf{w}} L_S(\mathbf{w}) = 0 \Leftrightarrow \mathbf{X}^\top \mathbf{X} \mathbf{w} = \mathbf{X}^\top \mathbf{y}$
- $\mathbf{X}^\top \mathbf{X}$ is a $(d+1) \times (d+1)$ matrix
- $\mathbf{X}^\top \mathbf{X}$ invertible: **Analytic solution** $\mathbf{w}_{\text{lin}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} = \mathbf{X}^\dagger \mathbf{y}$
- $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$: **pseudo-inverse** of \mathbf{X}
- **In practice**: use well-implemented \dagger routine for computing \mathbf{X}^\dagger

Hat matrix



- Approximate labels on inputs $\mathbf{x}_1, \dots, \mathbf{x}_m$:

$$\hat{\mathbf{y}} = \mathbf{X} \mathbf{w}_{\text{lin}} = \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} = \mathbf{H} \mathbf{y}$$

- $\mathbf{H} = \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ is called the **hat matrix** since it puts the “hat” on \mathbf{y}


Is linear regression really “learning”?

- No, in the sense that \mathbf{w}_{lin} has an **analytical solution**
- No algorithm necessary for iteratively improving the training loss
- Yes, in the sense that we achieve a small training loss $L_S(\mathbf{w})$
- Algorithm for computing the pseudo-inverse $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$

Linear classification vs. linear regression

- Minimizing $L_S(h) = \frac{1}{m} \sum_{i=1}^m \mathbb{I}[h(\mathbf{x}_i) \neq y_i]$ is **NP-hard**
- Minimizing $L_S(h) = \frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}_i) - y_i)^2$ has an **efficient analytical solution**
- **Idea:** $\{-1, +1\} \subset \mathbb{R} \Rightarrow$ use linear regression for classification!
- On input \mathbf{x} , predict label $\text{sign}(w_{\text{lin}}^\top \mathbf{x})$

Relationship between loss functions



- For label $+1$, square loss upper bounds 0-1 loss! (same for -1)
- Sacrifice **bound tightness** for **efficiency**

A second look at the Perceptron

What is the loss implicitly optimized by the PLA?

$$\begin{aligned}\mathbf{w}^{t+1} &\leftarrow \begin{cases} \mathbf{w}^t + y_i \mathbf{x}_i & \text{if } y_i \mathbf{w}^t \mathbf{x}_i < 0 \\ \mathbf{w}^t & \text{otherwise} \end{cases} \\ &\leftarrow \mathbf{w}^t - \nabla_{\mathbf{w}} \max(0, -y_i \mathbf{w}^t \mathbf{x}_i)\end{aligned}$$

A second look at the Perceptron

What is the loss implicitly optimized by the PLA?

$$\begin{aligned}\mathbf{w}^{t+1} &\leftarrow \begin{cases} \mathbf{w}^t + y_i \mathbf{x}_i & \text{if } y_i \mathbf{w}^t \mathbf{x}_i < 0 \\ \mathbf{w}^t & \text{otherwise} \end{cases} \\ &\leftarrow \mathbf{w}^t - \nabla_{\mathbf{w}} \max(0, -y_i \mathbf{w}^t \mathbf{x}_i)\end{aligned}$$

PLA follows the gradient of the local hinge loss

$$\ell_i(\mathbf{w}) = \max(0, -y_i \mathbf{w} \mathbf{x}_i)$$

Content

1 Supervised learning

2 Perceptron

3 Linear regression

4 Logistic regression


5 Non-linear transforms

6 Exercises

Logistic regression


- **Soft classification**: estimate **probability** of belonging to class
- Hypothesis $h(\mathbf{x}) = \theta(\sum_{i=0}^d w_i x_i) = \theta(\mathbf{w}^\top \mathbf{x})$
- **Logistic function**

$$\theta(s) = \frac{1}{1 + e^{-s}}$$



Logistic regression

- Assume that there are two classes $\{-1, +1\}$
- Ideally, data would be on the form $((\mathbf{x}_1, 0.8), \dots, (\mathbf{x}_m, 0.1))$, i.e. the **probability** of belonging to class $+1$
- However, data is usually on the form $((\mathbf{x}_1, +1), \dots, (\mathbf{x}_m, -1))$
- We can view labels as **hard probabilities**, i.e. 0 or 1



Logistic loss

- Find \mathbf{w} that minimizes

$$L_S(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \ell(h(\mathbf{x}_i), y_i)$$

- Cross-entropy loss $\ell(h(\mathbf{x}_i), y_i) = \ln(1 + \exp(-y_i \mathbf{w}^\top \mathbf{x}_i))$
- $L_S(\mathbf{w})$ is continuous, differentiable, and convex

Logistic loss

- Likelihood of a biased coin flip

$$p(y|\mathbf{x}, \mathbf{w}) = \theta(y\mathbf{w}^\top \mathbf{x})$$

Logistic loss

- Likelihood of a biased coin flip

$$p(y|\mathbf{x}, \mathbf{w}) = \theta(y\mathbf{w}^\top \mathbf{x})$$

- Logistic loss as logarithm of likelihood

$$\begin{aligned}\log p(y|\mathbf{x}, \mathbf{w}) &= \log \theta(y\mathbf{w}^\top \mathbf{x}) \\ &= \log \left(\frac{1}{1 + e^{-y\mathbf{w}^\top \mathbf{x}}} \right) \\ &= -\log(1 + e^{-y\mathbf{w}^\top \mathbf{x}})\end{aligned}$$

Logistic loss

- Likelihood of a biased coin flip

$$p(y|\mathbf{x}, \mathbf{w}) = \theta(y\mathbf{w}^\top \mathbf{x})$$

- Logistic loss as logarithm of likelihood

$$\begin{aligned}\log p(y|\mathbf{x}, \mathbf{w}) &= \log \theta(y\mathbf{w}^\top \mathbf{x}) \\ &= \log \left(\frac{1}{1 + e^{-y\mathbf{w}^\top \mathbf{x}}} \right) \\ &= -\log(1 + e^{-y\mathbf{w}^\top \mathbf{x}})\end{aligned}$$

- Probabilistic interpretation of losses → Bayesian formalism

Logistic loss

- Minimize $L_S(\mathbf{w}) \Leftrightarrow$ Find w such that $\nabla_{\mathbf{w}} L_S(\mathbf{w}) = 0$

- Gradient

$$\nabla_{\mathbf{w}} L_S(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \theta(-y_i \mathbf{w}^\top \mathbf{x}_i) (-y_i \mathbf{x}_i)$$


- Unfortunately, no analytic solution to $\nabla_{\mathbf{w}} L_S(\mathbf{w}) = 0$ descent)

Multiclass classification


- So far we have mainly talked about binary classification, i.e. $|\mathcal{Y}| = 2$
- What if there are more than two classes, i.e. $2 < |\mathcal{Y}| = k$?
- Two main approaches:
 - 1 **One-versus-all (OVA)**: **binary** classification of one class vs. rest
 - 2 **One-versus-one (OVO)**: **binary** classification of pairs of classes
- In both cases, perform soft binary classification (i.e. logistic regression) and return **most probable class**

One-versus-all

Binary classification:



Multi-class classification:



Multiclass classification

Properties of one-versus-all (OVA):

- Linear number of classifiers
- Imbalanced data!


Properties of one-versus-one (OVO):

- Quadratic number of classifiers
- More balanced data

Content


- 1 Supervised learning
- 2 Perceptron
- 3 Linear regression
- 4 Logistic regression
- 5 Non-linear transforms**
- 6 Exercises

Non-linear data



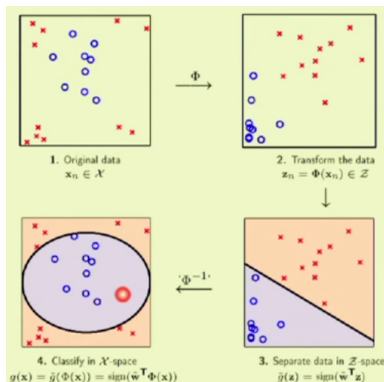
- Often data is not linearly separable at all
- **Idea**: derive **circular** perceptron, **circular** regression, etc.
- It would be better to take advantage of linear models!

Non-linear transforms



- Circular hypothesis: $h(\mathbf{x}) = \text{sign}(0.6 - x_1^2 - x_2^2)$
- Linear in **quadratic terms** x_1^2 and x_2^2 !
- Non-linear transform: introduce additional **non-linear terms**
- Quadratic transform: $\mathbf{z} = \Phi(\mathbf{x}) = (1, x_1, x_2, x_1x_2, x_1^2, x_2^2)$

Non-linear transforms



- Transform each data point (\mathbf{x}_i, y_i) to $(\mathbf{z}_i = \Phi(\mathbf{x}_i), y_i)$
- Apply linear algorithm in transformed space to find $\tilde{\mathbf{w}}$
- Hypothesis on input \mathbf{x} proportional to $\tilde{\mathbf{w}}^T \Phi(\mathbf{x})$

Price of non-linear transforms

- Let \mathcal{H}_q be the hypothesis class of q -th order polynomials
- Then $\mathcal{H}_1 \subset \mathcal{H}_2 \subset \dots \subset \mathcal{H}_q$
- However, number of features increases!
- The q -th order polynomial has $O(d^q)$ dimensions
- Increases memory requirements, running time of algorithms, etc.

Content

- 1 Supervised learning
- 2 Perceptron
- 3 Linear regression
- 4 Logistic regression
- 5 Non-linear transforms
- 6 Exercises**

Multiclass classification

- Assume that a binary classification algorithm \mathcal{A} runs in time N^3 on data of size N
- For a 10 class classification problem ($|\mathcal{Y}| = 10$), assume that there are exactly $N/10$ data points for each class
- What is the running time of OVA and OVO multiclass classification using algorithm \mathcal{A} ?

One-versus-all (OVA)

- We have to build 10 classifiers in total

One-versus-all (OVA)

- We have to build 10 classifiers in total
- Each classifier uses **all data**, i.e. N data points

One-versus-all (OVA)

- We have to build 10 classifiers in total
- Each classifier uses **all data**, i.e. N data points
- For each classifier, the running time of algorithm \mathcal{A} is N^3

One-versus-all (OVA)

- We have to build 10 classifiers in total
- Each classifier uses **all data**, i.e. N data points
- For each classifier, the running time of algorithm \mathcal{A} is N^3
- Hence the total running time is $10N^3$

One-versus-one (OVO)

- We have to build $\binom{10}{2} = \frac{10 \cdot 9}{2} = 45$ classifiers in total

One-versus-one (OVO)

- We have to build $\binom{10}{2} = \frac{10 \cdot 9}{2} = 45$ classifiers in total
- Each classifier uses only data points of 2 classes, i.e. $2N/10$

One-versus-one (OVO)

- We have to build $\binom{10}{2} = \frac{10 \cdot 9}{2} = 45$ classifiers in total
- Each classifier uses only data points of 2 classes, i.e. $2N/10$
- For each classifier, the running time of algorithm \mathcal{A} is $8N^3/1000$

One-versus-one (OVO)

- We have to build $\binom{10}{2} = \frac{10 \cdot 9}{2} = 45$ classifiers in total
- Each classifier uses only data points of 2 classes, i.e. $2N/10$
- For each classifier, the running time of algorithm \mathcal{A} is $8N^3/1000$
- Hence the total running time is $45 \cdot 8N^3/1000 = \frac{9}{25}N^3$

Hat matrix

Show that the hat matrix $\mathbf{H} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ has the following properties (where \mathbf{I} is the **identity matrix**):

1 \mathbf{H} is symmetric, i.e. $\mathbf{H}^\top = \mathbf{H}$

2 $\mathbf{H}^2 = \mathbf{H}$

3 $(\mathbf{I} - \mathbf{H})^2 = (\mathbf{I} - \mathbf{H})$

Hat matrix

Note that $(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top$ and that $(\mathbf{A}^{-1})^\top = (\mathbf{A}^\top)^{-1}$

$$\left\{ \mathbf{H}^\top = (\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top)^\top \right.$$

Hat matrix

Note that $(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top$ and that $(\mathbf{A}^{-1})^\top = (\mathbf{A}^\top)^{-1}$

$$\left\{ \begin{aligned} \mathbf{H}^\top &= (\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top)^\top \\ &= \mathbf{X} ((\mathbf{X}^\top \mathbf{X})^{-1})^\top \mathbf{X}^\top \end{aligned} \right.$$

Hat matrix

Note that $(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top$ and that $(\mathbf{A}^{-1})^\top = (\mathbf{A}^\top)^{-1}$

$$\left\{ \begin{aligned} \mathbf{H}^\top &= (\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top)^\top \\ &= \mathbf{X} ((\mathbf{X}^\top \mathbf{X})^{-1})^\top \mathbf{X}^\top \\ &= \mathbf{X} ((\mathbf{X}^\top \mathbf{X})^\top)^{-1} \mathbf{X}^\top \end{aligned} \right.$$

Hat matrix

Note that $(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top$ and that $(\mathbf{A}^{-1})^\top = (\mathbf{A}^\top)^{-1}$

$$\left\{ \begin{aligned} \mathbf{H}^\top &= (\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top)^\top \\ &= \mathbf{X} ((\mathbf{X}^\top \mathbf{X})^{-1})^\top \mathbf{X}^\top \\ &= \mathbf{X} ((\mathbf{X}^\top \mathbf{X})^\top)^{-1} \mathbf{X}^\top \\ &= \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top = \mathbf{H} \end{aligned} \right.$$

Hat matrix

Note that $\mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$ and $\mathbf{A}\mathbf{I} = \mathbf{I}\mathbf{A} = \mathbf{A}$

$$\left\{ \begin{aligned} \mathbf{H}^2 &= \mathbf{H} \cdot \mathbf{H} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \end{aligned} \right.$$

Hat matrix

Note that $\mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$ and $\mathbf{A}\mathbf{I} = \mathbf{I}\mathbf{A} = \mathbf{A}$

$$\left\{ \begin{aligned} \mathbf{H}^2 &= \mathbf{H} \cdot \mathbf{H} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \\ &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{X}) (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \end{aligned} \right.$$

Hat matrix

Note that $\mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$ and $\mathbf{A}\mathbf{I} = \mathbf{I}\mathbf{A} = \mathbf{A}$

$$\left\{ \begin{aligned} \mathbf{H}^2 &= \mathbf{H} \cdot \mathbf{H} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \\ &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{X}) (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \\ &= \mathbf{X} \mathbf{I} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \end{aligned} \right.$$

Hat matrix

Note that $\mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$ and $\mathbf{A}\mathbf{I} = \mathbf{I}\mathbf{A} = \mathbf{A}$

$$\left\{ \begin{aligned} \mathbf{H}^2 &= \mathbf{H} \cdot \mathbf{H} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \\ &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{X}) (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \\ &= \mathbf{X} \mathbf{I} (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \\ &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top = \mathbf{H} \end{aligned} \right.$$

Hat matrix

$$\left\{ \begin{aligned} (\mathbf{I} - \mathbf{H})^2 &= \mathbf{I}^2 - 2\mathbf{I}\mathbf{H} + \mathbf{H}^2 \end{aligned} \right.$$

Hat matrix

$$\begin{cases} (\mathbf{I} - \mathbf{H})^2 = \mathbf{I}^2 - 2\mathbf{I}\mathbf{H} + \mathbf{H}^2 \\ \quad \quad \quad = \mathbf{I} - 2\mathbf{H} + \mathbf{H} = \mathbf{I} - \mathbf{H} \end{cases}$$

Perceptron learning algorithm (PLA)

Perceptron learning algorithm

- 1 Initialize weight vector $\mathbf{w}^0 = 0$
- 2 Find a **mistake** (\mathbf{x}_i, y_i) such that $h(\mathbf{x}_i) \neq y_i$
- 3 Update weights as $\mathbf{w}_1 \leftarrow \mathbf{w}^0 + y_i \mathbf{x}_i$
- 4 Repeat from 2. for weight vector \mathbf{w}^t , $t = 1, 2, \dots$

Letting $R^2 = \max_i \|\mathbf{x}_i\|^2$, show that after t iterations, $\|\mathbf{w}^t\|^2 \leq tR^2$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Consider a single step of PLA:

$$\left\{ \begin{array}{l} \|\mathbf{w}^{t+1}\|^2 = \|\mathbf{w}^t + y_i \mathbf{x}_i\|^2 \\ \end{array} \right.$$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Consider a single step of PLA:

$$\left\{ \begin{aligned} \|\mathbf{w}^{t+1}\|^2 &= \|\mathbf{w}^t + y_i \mathbf{x}_i\|^2 \\ &= \|\mathbf{w}^t\|^2 + 2y_i \mathbf{w}^{t\top} \mathbf{x}_i + y_i^2 \|\mathbf{x}_i\|^2 \end{aligned} \right.$$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Consider a single step of PLA:

$$\left\{ \begin{aligned} \|\mathbf{w}^{t+1}\|^2 &= \|\mathbf{w}^t + y_i \mathbf{x}_i\|^2 \\ &= \|\mathbf{w}^t\|^2 + 2y_i \mathbf{w}^{t\top} \mathbf{x}_i + y_i^2 \|\mathbf{x}_i\|^2 \\ &\leq \|\mathbf{w}^t\|^2 + 0 + \|\mathbf{x}_i\|^2 \end{aligned} \right.$$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Consider a single step of PLA:

$$\left\{ \begin{array}{l} \|\mathbf{w}^{t+1}\|^2 = \|\mathbf{w}^t + y_i \mathbf{x}_i\|^2 \\ \quad = \|\mathbf{w}^t\|^2 + 2y_i \mathbf{w}^{t\top} \mathbf{x}_i + y_i^2 \|\mathbf{x}_i\|^2 \\ \leq \|\mathbf{w}^t\|^2 + 0 + \|\mathbf{x}_i\|^2 \\ \leq \|\mathbf{w}^t\|^2 + R^2 \end{array} \right.$$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Consider a single step of PLA:

$$\left\{ \begin{array}{l} \|\mathbf{w}^{t+1}\|^2 = \|\mathbf{w}^t + y_i \mathbf{x}_i\|^2 \\ \quad = \|\mathbf{w}^t\|^2 + 2y_i \mathbf{w}^{t\top} \mathbf{x}_i + y_i^2 \|\mathbf{x}_i\|^2 \\ \quad \leq \|\mathbf{w}^t\|^2 + 0 + \|\mathbf{x}_i\|^2 \\ \quad \leq \|\mathbf{w}^t\|^2 + R^2 \end{array} \right.$$

Perceptron learning algorithm (PLA)

An update happens when prediction are wrong

$$\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i \equiv y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$$

Consider a single step of PLA:

$$\left\{ \begin{aligned} \|\mathbf{w}^{t+1}\|^2 &= \|\mathbf{w}^t + y_i \mathbf{x}_i\|^2 \\ &= \|\mathbf{w}^t\|^2 + 2y_i \mathbf{w}^{t\top} \mathbf{x}_i + y_i^2 \|\mathbf{x}_i\|^2 \\ &\leq \|\mathbf{w}^t\|^2 + 0 + \|\mathbf{x}_i\|^2 \\ &\leq \|\mathbf{w}^t\|^2 + R^2 \end{aligned} \right.$$

Hence after t iterations, $\|\mathbf{w}^t\|^2 \leq \|\mathbf{w}^0\|^2 + tR^2 = 0 + tR^2 = tR^2$

Perceptron learning algorithm (PLA)

Data linearly separable \Rightarrow exists \mathbf{w}_* such that $y_i \mathbf{w}_*^\top \mathbf{x}_i > 0, \forall i \in [m]$

Perceptron learning algorithm

- 1 Initialize weight vector $\mathbf{w}^0 = \mathbf{0}$
- 2 Find a **mistake** (\mathbf{x}_i, y_i) such that $h(\mathbf{x}_i) \neq y_i$
- 3 Update weights as $\mathbf{w}^1 \leftarrow \mathbf{w}^0 + y_i \mathbf{x}_i$
- 4 Repeat from 2. for weight vector $\mathbf{w}^t, t = 1, 2, \dots$

Letting $\rho = \min_i y_i \frac{\mathbf{w}_*^\top \mathbf{x}_i}{\|\mathbf{w}_*\|} > 0$, show that after t iterations, $\frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\|} \geq t\rho$

Perceptron learning algorithm (PLA)

Consider a single step of PLA:

$$\left\{ \frac{\mathbf{w}_*^T \mathbf{w}^{t+1}}{\|\mathbf{w}_*\|} = \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + y_i \frac{\mathbf{w}_*^T}{\|\mathbf{w}_*\|} \mathbf{x}_i \right.$$

Perceptron learning algorithm (PLA)

Consider a single step of PLA:

$$\left\{ \begin{array}{l} \frac{\mathbf{w}_*^T \mathbf{w}^{t+1}}{\|\mathbf{w}_*\|} = \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + y_i \frac{\mathbf{w}_*^T}{\|\mathbf{w}_*\|} \mathbf{x}_i \\ \geq \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + \rho \end{array} \right.$$

Perceptron learning algorithm (PLA)

Consider a single step of PLA:

$$\left\{ \begin{array}{l} \frac{\mathbf{w}_*^T \mathbf{w}^{t+1}}{\|\mathbf{w}_*\|} = \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + y_i \frac{\mathbf{w}_*^T}{\|\mathbf{w}_*\|} \mathbf{x}_i \\ \geq \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + \rho \end{array} \right.$$

Perceptron learning algorithm (PLA)

Consider a single step of PLA:

$$\begin{cases} \frac{\mathbf{w}_*^T \mathbf{w}^{t+1}}{\|\mathbf{w}_*\|} = \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + y_i \frac{\mathbf{w}_*^T}{\|\mathbf{w}_*\|} \mathbf{x}_i \\ \geq \frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} + \rho \end{cases}$$

Hence after t iterations, $\frac{\mathbf{w}_*^T \mathbf{w}^t}{\|\mathbf{w}_*\|} \geq \frac{\mathbf{w}_*^T \mathbf{w}_0}{\|\mathbf{w}_*\|} + t\rho = 0 + t\rho = t\rho$

Perceptron learning algorithm (PLA)

- Putting the two results together, we get

$$\frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \frac{t\rho}{\sqrt{t}R} = \sqrt{t} \frac{\rho}{R}$$

Perceptron learning algorithm (PLA)

- Putting the two results together, we get

$$\frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \frac{t\rho}{\sqrt{t}R} = \sqrt{t} \frac{\rho}{R}$$

- The scalar product of two unit vectors is upper bounded by 1:

$$1 \geq \frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \sqrt{t} \frac{\rho}{R}$$

Perceptron learning algorithm (PLA)

- Putting the two results together, we get

$$\frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \frac{t\rho}{\sqrt{t}R} = \sqrt{t} \frac{\rho}{R}$$

- The scalar product of two unit vectors is upper bounded by 1:

$$1 \geq \frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \sqrt{t} \frac{\rho}{R}$$

- It follows that

$$\sqrt{t} \leq \frac{R}{\rho} \Leftrightarrow t \leq \frac{R^2}{\rho^2}$$

Perceptron learning algorithm (PLA)

- Putting the two results together, we get

$$\frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \frac{t\rho}{\sqrt{t}R} = \sqrt{t} \frac{\rho}{R}$$

- The scalar product of two unit vectors is upper bounded by 1:

$$1 \geq \frac{\mathbf{w}_*^\top \mathbf{w}^t}{\|\mathbf{w}_*\| \|\mathbf{w}^t\|} \geq \sqrt{t} \frac{\rho}{R}$$

- It follows that

$$\sqrt{t} \leq \frac{R}{\rho} \Leftrightarrow t \leq \frac{R^2}{\rho^2}$$

- Hence PLA converges after at most R^2/ρ^2 iterations!